

exploration_template

April 20, 2021

1 2019.02 Ford Go Bike - tripdata

1.1 by Szymon Debski

1.2 Preliminary Wrangling

This data set is taken from <https://www.fordgobike.com/system-data> and represents trips taken by members of the service for month of February of 2019.

Data consists of info about trips taken by service's members, their types, their age, their gender, stations of starting and ending trips, duration of trips etc.

```
[147]: # import all packages and set plots to be embedded inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb

%matplotlib inline

import warnings
warnings.simplefilter("ignore")
```

Load in your dataset and describe its properties through the questions below. Try and motivate your exploration goals through this section.

```
[148]: df = pd.read_csv('201902-fordgobike-tripdata.csv')
df.head()
```

```
[148]:
```

	duration_sec	start_time	end_time	\
0	52185	2019-02-28 17:32:10.1450	2019-03-01 08:01:55.9750	
1	42521	2019-02-28 18:53:21.7890	2019-03-01 06:42:03.0560	
2	61854	2019-02-28 12:13:13.2180	2019-03-01 05:24:08.1460	
3	36490	2019-02-28 17:54:26.0100	2019-03-01 04:02:36.8420	
4	1585	2019-02-28 23:54:18.5490	2019-03-01 00:20:44.0740	

	start_station_id	start_station_name	\
0	21.0	Montgomery St BART Station (Market St at 2nd St)	
1	23.0	The Embarcadero at Steuart St	
2	86.0	Market St at Dolores St	

3	375.0	Grove St at Masonic Ave
4	7.0	Frank H Ogawa Plaza

	start_station_latitude	start_station_longitude	end_station_id \
0	37.789625	-122.400811	13.0
1	37.791464	-122.391034	81.0
2	37.769305	-122.426826	3.0
3	37.774836	-122.446546	70.0
4	37.804562	-122.271738	222.0

	end_station_name	end_station_latitude \
0	Commercial St at Montgomery St	37.794231
1	Berry St at 4th St	37.775880
2	Powell St BART Station (Market St at 4th St)	37.786375
3	Central Ave at Fell St	37.773311
4	10th Ave at E 15th St	37.792714

	end_station_longitude	bike_id	user_type	member_birth_year \
0	-122.402923	4902	Customer	1984.0
1	-122.393170	2535	Customer	NaN
2	-122.404904	5905	Customer	1972.0
3	-122.444293	6638	Subscriber	1989.0
4	-122.248780	4898	Subscriber	1974.0

	member_gender	bike_share_for_all_trip
0	Male	No
1	NaN	No
2	Male	No
3	Other	No
4	Male	Yes

```
[149]: df.shape
```

```
[149]: (183412, 16)
```

```
[150]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 183412 entries, 0 to 183411
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   duration_sec                          183412 non-null int64
1   start_time                            183412 non-null object
2   end_time                              183412 non-null object
3   start_station_id                      183215 non-null float64
4   start_station_name                    183215 non-null object
5   start_station_latitude                183412 non-null float64
```

```

6  start_station_longitude 183412 non-null float64
7  end_station_id          183215 non-null float64
8  end_station_name        183215 non-null object
9  end_station_latitude    183412 non-null float64
10 end_station_longitude   183412 non-null float64
11 bike_id                 183412 non-null int64
12 user_type               183412 non-null object
13 member_birth_year       175147 non-null float64
14 member_gender           175147 non-null object
15 bike_share_for_all_trip 183412 non-null object
dtypes: float64(7), int64(2), object(7)
memory usage: 22.4+ MB

```

```
[151]: df.describe()
```

```

[151]:
      duration_sec  start_station_id  start_station_latitude \
count  183412.000000      183215.000000      183412.000000
mean      726.078435      138.590427      37.771223
std      1794.389780      111.778864      0.099581
min        61.000000       3.000000      37.317298
25%       325.000000      47.000000      37.770083
50%       514.000000     104.000000      37.780760
75%       796.000000     239.000000      37.797280
max      85444.000000     398.000000      37.880222

      start_station_longitude  end_station_id  end_station_latitude \
count      183412.000000      183215.000000      183412.000000
mean        -122.352664      136.249123      37.771427
std           0.117097      111.515131      0.099490
min        -122.453704       3.000000      37.317298
25%        -122.412408      44.000000      37.770407
50%        -122.398285     100.000000      37.781010
75%        -122.286533     235.000000      37.797320
max        -121.874119     398.000000      37.880222

      end_station_longitude      bike_id  member_birth_year
count      183412.000000  183412.000000      175147.000000
mean        -122.352250    4472.906375      1984.806437
std           0.116673    1664.383394      10.116689
min        -122.453704      11.000000      1878.000000
25%        -122.411726    3777.000000      1980.000000
50%        -122.398279    4958.000000      1987.000000
75%        -122.288045    5502.000000      1992.000000
max        -121.874119    6645.000000      2001.000000

```

1.2.1 Data Cleaning

```
[152]: df_clean = df.copy()
```

Remove columns that are not of interest to me

```
[153]: list(df_clean.columns)
```

```
[153]: ['duration_sec',
        'start_time',
        'end_time',
        'start_station_id',
        'start_station_name',
        'start_station_latitude',
        'start_station_longitude',
        'end_station_id',
        'end_station_name',
        'end_station_latitude',
        'end_station_longitude',
        'bike_id',
        'user_type',
        'member_birth_year',
        'member_gender',
        'bike_share_for_all_trip']
```

```
[154]: df_clean.drop(['start_station_id',
        'start_station_name',
        'start_station_latitude',
        'start_station_longitude',
        'end_station_id',
        'end_station_name',
        'end_station_latitude',
        'end_station_longitude'], axis=1, inplace=True)
```

```
[155]: df_clean.head()
```

```
[155]:
```

	duration_sec		start_time		end_time	bike_id \
0	52185	2019-02-28	17:32:10.1450	2019-03-01	08:01:55.9750	4902
1	42521	2019-02-28	18:53:21.7890	2019-03-01	06:42:03.0560	2535
2	61854	2019-02-28	12:13:13.2180	2019-03-01	05:24:08.1460	5905
3	36490	2019-02-28	17:54:26.0100	2019-03-01	04:02:36.8420	6638
4	1585	2019-02-28	23:54:18.5490	2019-03-01	00:20:44.0740	4898

	user_type	member_birth_year	member_gender	bike_share_for_all_trip
0	Customer	1984.0	Male	No
1	Customer	NaN	NaN	No
2	Customer	1972.0	Male	No
3	Subscriber	1989.0	Other	No
4	Subscriber	1974.0	Male	Yes

```
[156]: df_clean.isnull().sum()
```

```
[156]: duration_sec          0
start_time                0
end_time                  0
bike_id                   0
user_type                 0
member_birth_year        8265
member_gender             8265
bike_share_for_all_trip   0
dtype: int64
```

Drop records with null values

```
[157]: df_clean = df_clean[df_clean['member_birth_year'].isnull() == False]
```

```
[158]: df_clean.isnull().sum()
```

```
[158]: duration_sec          0
start_time                0
end_time                  0
bike_id                   0
user_type                 0
member_birth_year        0
member_gender             0
bike_share_for_all_trip   0
dtype: int64
```

```
[159]: df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 175147 entries, 0 to 183411
Data columns (total 8 columns):
#   Column              Non-Null Count  Dtype
---  -
0   duration_sec        175147 non-null  int64
1   start_time          175147 non-null  object
2   end_time            175147 non-null  object
3   bike_id             175147 non-null  int64
4   user_type           175147 non-null  object
5   member_birth_year   175147 non-null  float64
6   member_gender       175147 non-null  object
7   bike_share_for_all_trip 175147 non-null  object
dtypes: float64(1), int64(2), object(5)
memory usage: 12.0+ MB
```

```
[160]: df_clean.user_type.value_counts()
```

```
[160]: Subscriber    158516
      Customer      16631
      Name: user_type, dtype: int64
```

```
[161]: df_clean.member_gender.value_counts()
```

```
[161]: Male        130651
      Female     40844
      Other       3652
      Name: member_gender, dtype: int64
```

Drop 'Other' from gender column'

```
[162]: df_clean = df_clean[df_clean['member_gender'] != 'Other']
```

```
[163]: df_clean.member_gender.value_counts()
```

```
[163]: Male        130651
      Female     40844
      Name: member_gender, dtype: int64
```

```
[164]: df_clean.duplicated().sum()
```

```
[164]: 0
```

```
[165]: df_clean.start_time = pd.to_datetime(df_clean['start_time'])
      df_clean.end_time = pd.to_datetime(df_clean['end_time'])
```

```
[166]: df_clean['weekday'] = df_clean['start_time'].dt.day_name()
      df_clean['hour'] = df_clean['start_time'].dt.hour
```

```
[167]: df_clean.sample(1)
```

```
[167]:      duration_sec      start_time      end_time  bike_id \
28155          359 2019-02-25 08:57:11.110 2019-02-25 09:03:10.381    5633

      user_type  member_birth_year  member_gender  bike_share_for_all_trip \
28155  Subscriber          1987.0        Female                No

      weekday  hour
28155  Monday      8
```

```
[168]: df_clean.weekday.value_counts()
```

```
[168]: Thursday     33005
      Tuesday    30051
      Wednesday  27853
      Friday     27102
      Monday     25151
```

```
Sunday      14203
Saturday    14130
Name: weekday, dtype: int64
```

Create a column duration in minutes

```
[171]: df_clean['duration_minutes'] = round(df_clean['duration_sec'] / 60, 0)
```

```
[172]: df_clean.sample(5)
```

```
[172]:
```

	duration_sec	start_time	end_time	bike_id	\
182441	714	2019-02-01 08:11:47.664	2019-02-01 08:23:41.687	1469	
34008	306	2019-02-24 09:48:33.715	2019-02-24 09:53:40.600	5940	
142935	285	2019-02-07 16:32:46.908	2019-02-07 16:37:32.328	2848	
17846	456	2019-02-26 20:23:55.220	2019-02-26 20:31:31.586	1576	
121410	574	2019-02-11 15:19:28.770	2019-02-11 15:29:03.413	4895	

	user_type	member_birth_year	member_gender	bike_share_for_all_trip	\
182441	Subscriber	1992.0	Female	No	
34008	Subscriber	1987.0	Male	No	
142935	Subscriber	1982.0	Male	No	
17846	Customer	1991.0	Male	No	
121410	Customer	1997.0	Male	No	

	weekday	hour	duration_minutes
182441	Friday	8	12.0
34008	Sunday	9	5.0
142935	Thursday	16	5.0
17846	Tuesday	20	8.0
121410	Monday	15	10.0

Calculate age for users of service

```
[173]: df_clean['age'] = 2019 - df_clean['member_birth_year']
```

```
[174]: df_clean.sample(5)
```

```
[174]:
```

	duration_sec	start_time	end_time	bike_id	\
78865	175	2019-02-18 18:34:18.671	2019-02-18 18:37:13.813	4909	
33702	625	2019-02-24 10:50:43.535	2019-02-24 11:01:09.124	4065	
25297	883	2019-02-25 17:04:49.919	2019-02-25 17:19:33.485	6353	
108268	534	2019-02-13 08:57:10.372	2019-02-13 09:06:04.484	3676	
157698	579	2019-02-05 20:54:28.249	2019-02-05 21:04:07.592	2859	

	user_type	member_birth_year	member_gender	bike_share_for_all_trip	\
78865	Subscriber	1983.0	Male	No	
33702	Subscriber	1995.0	Male	Yes	
25297	Subscriber	1990.0	Male	No	
108268	Customer	1989.0	Male	No	

157698	Subscriber		1994.0	Male	Yes
--------	------------	--	--------	------	-----

	weekday	hour	duration_minutes	age
78865	Monday	18	3.0	36.0
33702	Sunday	10	10.0	24.0
25297	Monday	17	15.0	29.0
108268	Wednesday	8	9.0	30.0
157698	Tuesday	20	10.0	25.0

```
[175]: df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 171495 entries, 0 to 183411
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   duration_sec                          171495 non-null  int64
1   start_time                            171495 non-null  datetime64[ns]
2   end_time                              171495 non-null  datetime64[ns]
3   bike_id                               171495 non-null  int64
4   user_type                             171495 non-null  object
5   member_birth_year                     171495 non-null  float64
6   member_gender                          171495 non-null  object
7   bike_share_for_all_trip                171495 non-null  object
8   weekday                               171495 non-null  object
9   hour                                   171495 non-null  int64
10  duration_minutes                       171495 non-null  float64
11  age                                    171495 non-null  float64
dtypes: datetime64[ns](2), float64(3), int64(3), object(4)
memory usage: 17.0+ MB
```

```
[176]: df_clean['duration_minutes'] = df_clean['duration_minutes'].astype(int)
df_clean['age'] = df_clean['age'].astype(int)
```

```
[177]: df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 171495 entries, 0 to 183411
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   duration_sec                          171495 non-null  int64
1   start_time                            171495 non-null  datetime64[ns]
2   end_time                              171495 non-null  datetime64[ns]
3   bike_id                               171495 non-null  int64
4   user_type                             171495 non-null  object
5   member_birth_year                     171495 non-null  float64
6   member_gender                          171495 non-null  object
```



```

7  bike_share_for_all_trip  171495 non-null  object
8  weekday                 171495 non-null  object
9  hour                    171495 non-null  int64
10 duration_minutes        171495 non-null  int32
11 age                     171495 non-null  int32
dtypes: datetime64[ns](2), float64(1), int32(2), int64(3), object(4)
memory usage: 15.7+ MB

```

1.2.2 What is the structure of your dataset?

The data set consists of approx. 180k bike rides in San Francisco.

The data contains:

- Trip duration
- Start / Endstation
- Bike ID
- User info (user type, date of birth, gender)

1.2.3 What is the main feature of interest in your dataset?

I'm interested in analyzing trip duration depending on user type

1.2.4 What features in the dataset do you think will help support your investigation into your feature(s) of interest?

- duration_sec
- user_type
- member_birth_year
- member_gender

1.3 Univariate Exploration

In this section, investigate distributions of individual variables. If you see unusual points or outliers, take a deeper look to clean things up and prepare yourself to look at relationships between variables.

```
[178]: df_clean.describe()
```

```

[178]:
      duration_sec  bike_id  member_birth_year  hour \
count  171495.000000  171495.000000  171495.000000  171495.000000
mean      697.987218    4480.962868    1984.842328    13.451698
std      1576.717221    1658.635854     10.113921     4.732831
min        61.000000     11.000000    1878.000000     0.000000
25%       323.000000    3799.000000    1980.000000     9.000000
50%       510.000000    4959.000000    1987.000000    14.000000
75%       787.000000    5505.000000    1992.000000    17.000000
max      84548.000000    6645.000000    2001.000000    23.000000


      duration_minutes  age

```

count	171495.000000	171495.000000
mean	11.632992	34.157672
std	26.280664	10.113921
min	1.000000	18.000000
25%	5.000000	27.000000
50%	8.000000	32.000000
75%	13.000000	39.000000
max	1409.000000	141.000000

It seems that the 'duration_minutes' data is highly skewed to the right, most rides are short however there are some that are long we will have to take a closer look

```
[179]: bins = np.arange(0, df_clean['duration_minutes'].max()+10, 10)
plt.hist(df_clean['duration_minutes'], bins = bins);
```



```
[180]: np.log10(df_clean['duration_minutes'].describe())
```

```
[180]: count    5.234251
mean      1.065691
std       1.419636
min       0.000000
25%      0.698970
50%      0.903090
75%      1.113943
max       3.148911
Name: duration_minutes, dtype: float64
```


Let's transform the scale so that we have a better look at the data. We will use the log10 transformation for the bins and use a log scale for the x-axis.

```
[181]: bin_size = 0.075
bins = 10 ** np.arange(0, np.log10(df_clean['duration_minutes'].max()) + 1,
    ↪ bin_size, bin_size)

plt.figure(figsize=[10, 7]);
plt.hist(data = df_clean, x = 'duration_minutes', bins = bins);
plt.xscale('log');
ticks = [1, 3, 5, 10, 30, 100, 300, 1000, 2000]
labels = ['{}'.format(val) for val in ticks]

plt.xticks(ticks, labels)
plt.xlabel('Duration')
```

```
[181]: Text(0.5, 0, 'Duration')
```



This looks much better. We can drop the outliers - trips longer than 100 minutes.

```
[182]: df_clean = df_clean[df_clean['duration_minutes'] <= 100]
df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 170834 entries, 4 to 183411
```

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	duration_sec	170834 non-null	int64
1	start_time	170834 non-null	datetime64[ns]
2	end_time	170834 non-null	datetime64[ns]
3	bike_id	170834 non-null	int64
4	user_type	170834 non-null	object
5	member_birth_year	170834 non-null	float64
6	member_gender	170834 non-null	object
7	bike_share_for_all_trip	170834 non-null	object
8	weekday	170834 non-null	object
9	hour	170834 non-null	int64
10	duration_minutes	170834 non-null	int32
11	age	170834 non-null	int32


dtypes: datetime64[ns](2), float64(1), int32(2), int64(3), object(4)
memory usage: 15.6+ MB

```
[183]: bin_size = 0.075
bins = 10 ** np.arange(0, np.log10(df_clean['duration_minutes'].max()) + 1,
    ↪ bin_size, bin_size)

plt.figure(figsize=[10, 7]);
plt.hist(data = df_clean, x = 'duration_minutes', bins = bins);
plt.xscale('log');
ticks = [1, 3, 5, 10, 30, 100, 300, 1000, 2000]
labels = ['{}'.format(val) for val in ticks]

plt.xticks(ticks, labels)
plt.xlabel('Duration')
```

```
[183]: Text(0.5, 0, 'Duration')
```




```
[184]: df_clean['age'].describe()
```

```
[184]: count    170834.000000
      mean      34.153084
      std      10.113001
      min      18.000000
      25%      27.000000
      50%      32.000000
      75%      39.000000
      max      141.000000
      Name: age, dtype: float64
```

Data skewed to the right plus oldest customer is 141 years old this can't be correct

```
[185]: bins = np.arange(18, df_clean['age'].max()+2, 2)
      plt.hist(df_clean['age'], bins = bins);
      plt.xlabel('Age');
```



The distribution does not look good, let's drop the outliers

```
[186]: df_clean = df_clean[df_clean['age'] <= 65]
df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 169513 entries, 4 to 183411
```

```
Data columns (total 12 columns):
```


#	Column	Non-Null Count	Dtype
0	duration_sec	169513 non-null	int64
1	start_time	169513 non-null	datetime64[ns]
2	end_time	169513 non-null	datetime64[ns]
3	bike_id	169513 non-null	int64
4	user_type	169513 non-null	object
5	member_birth_year	169513 non-null	float64
6	member_gender	169513 non-null	object
7	bike_share_for_all_trip	169513 non-null	object
8	weekday	169513 non-null	object
9	hour	169513 non-null	int64
10	duration_minutes	169513 non-null	int32
11	age	169513 non-null	int32

```
dtypes: datetime64[ns](2), float64(1), int32(2), int64(3), object(4)
```

```
memory usage: 15.5+ MB
```

This looks much better still right-skewed but in reason

```
[187]: bins = np.arange(18, df_clean['age'].max()+1, 1)
plt.hist(df_clean['age'], bins = bins);
plt.xlabel('Age');
```



```
[188]: df_clean.head()
```

```
[188]:
```

	duration_sec	start_time	end_time	bike_id	\
4	1585	2019-02-28 23:54:18.549	2019-03-01 00:20:44.074	4898	
5	1793	2019-02-28 23:49:58.632	2019-03-01 00:19:51.760	5200	
6	1147	2019-02-28 23:55:35.104	2019-03-01 00:14:42.588	3803	
7	1615	2019-02-28 23:41:06.766	2019-03-01 00:08:02.756	6329	
9	1049	2019-02-28 23:49:47.699	2019-03-01 00:07:17.025	6488	


	user_type	member_birth_year	member_gender	bike_share_for_all_trip	\
4	Subscriber	1974.0	Male	Yes	
5	Subscriber	1959.0	Male	No	
6	Subscriber	1983.0	Female	No	
7	Subscriber	1989.0	Male	No	
9	Subscriber	1992.0	Male	No	

	weekday	hour	duration_minutes	age
4	Thursday	23	26	45
5	Thursday	23	30	60
6	Thursday	23	19	36

7	Thursday	23	27	30
9	Thursday	23	17	27


When looking at our customer base we can see that most people which use the bikes are subscribers

```
[189]: sb.countplot(data = df_clean, x = 'user_type', color = sb.color_palette()[0])  
plt.xticks(rotation = 45);
```



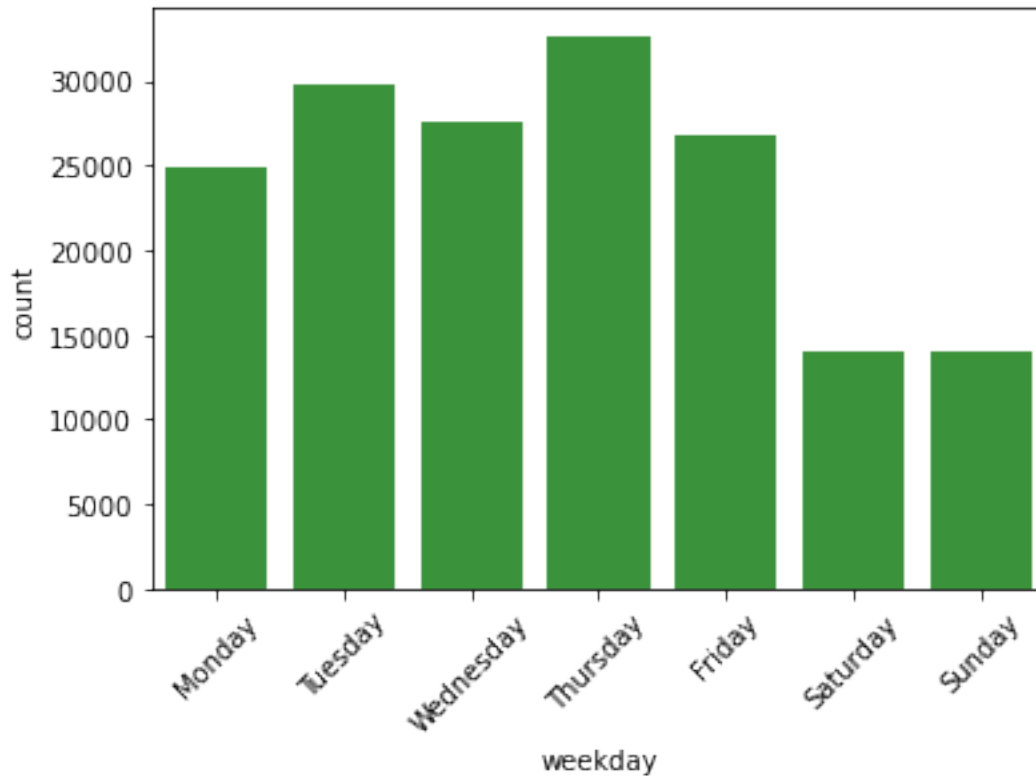
Most people using the service are Male

```
[190]: sb.countplot(data = df_clean, x = 'member_gender', color = sb.  
    ↳ color_palette()[1])  
plt.xticks(rotation = 45);
```

The highest numbers of rides occurred on Thursday

```
[191]: days = [ 'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday',  
               ↪ 'Sunday']  
  
sb.countplot(data = df_clean, x = 'weekday', color = sb.color_palette()[2],  
             ↪ order = days)  
plt.xticks(rotation = 45);
```



Make sure that, after every plot or related series of plots, that you include a Markdown cell with comments about what you observed, and what you plan on investigating next.

1.3.1 Discuss the distribution(s) of your variable(s) of interest. Were there any unusual points? Did you need to perform any transformations?

First we analyzed the 'duration in minutes' as it is our main variable of interest. It was heavily right-skewed. We had to transform the data using a log scale which helped us get a normal distribution. We also dropped trips longer than 100 minutes which made the data even cleaner. At first glance, we can see that most trips were around 10 minutes long.

1.3.2 Of the features you investigated, were there any unusual distributions? Did you perform any operations on the data to tidy, adjust, or change the form of the data? If so, why did you do this?

The distribution of the 'age' variable was also heavily right-skewed. When looking closely at the age data there were some unusual data points - the oldest customer was 141 years old - which for sure is a mistake. For our analysis, I decided to drop customers older than 65 years old. This in turn made our distribution nicer - still right-skewed but in reason.

1.4 Bivariate Exploration

In this section, investigate relationships between pairs of variables in your data. Make sure the variables that you cover here have been introduced in some fashion in the previous section (univariate exploration).

```
[192]: df_clean.describe()
```


```
[192]:
```

	duration_sec	bike_id	member_birth_year	hour \
count	169513.000000	169513.000000	169513.000000	169513.000000
mean	632.194392	4481.961047	1985.156489	13.457404
std	504.899629	1657.785239	9.460529	4.738949
min	61.000000	11.000000	1954.000000	0.000000
25%	322.000000	3802.000000	1980.000000	9.000000
50%	508.000000	4959.000000	1987.000000	14.000000
75%	782.000000	5504.000000	1992.000000	17.000000
max	6020.000000	6645.000000	2001.000000	23.000000

	duration_minutes	age
count	169513.000000	169513.000000
mean	10.536478	33.843511
std	8.418791	9.460529
min	1.000000	18.000000
25%	5.000000	27.000000
50%	8.000000	32.000000
75%	13.000000	39.000000
max	100.000000	65.000000


We can see a negative relationship between age and the duration of the ride - which is understandable - however, this does not mean that older people only ride for a short time there are many cases where over 65 rode for over 80 minutes.

```
[193]: plt.figure(figsize = [8, 6])
plt.scatter(data = df_clean, x = 'age', y = 'duration_minutes', alpha = 1/10);
plt.xlabel('Age');
plt.ylabel('Duration');
```



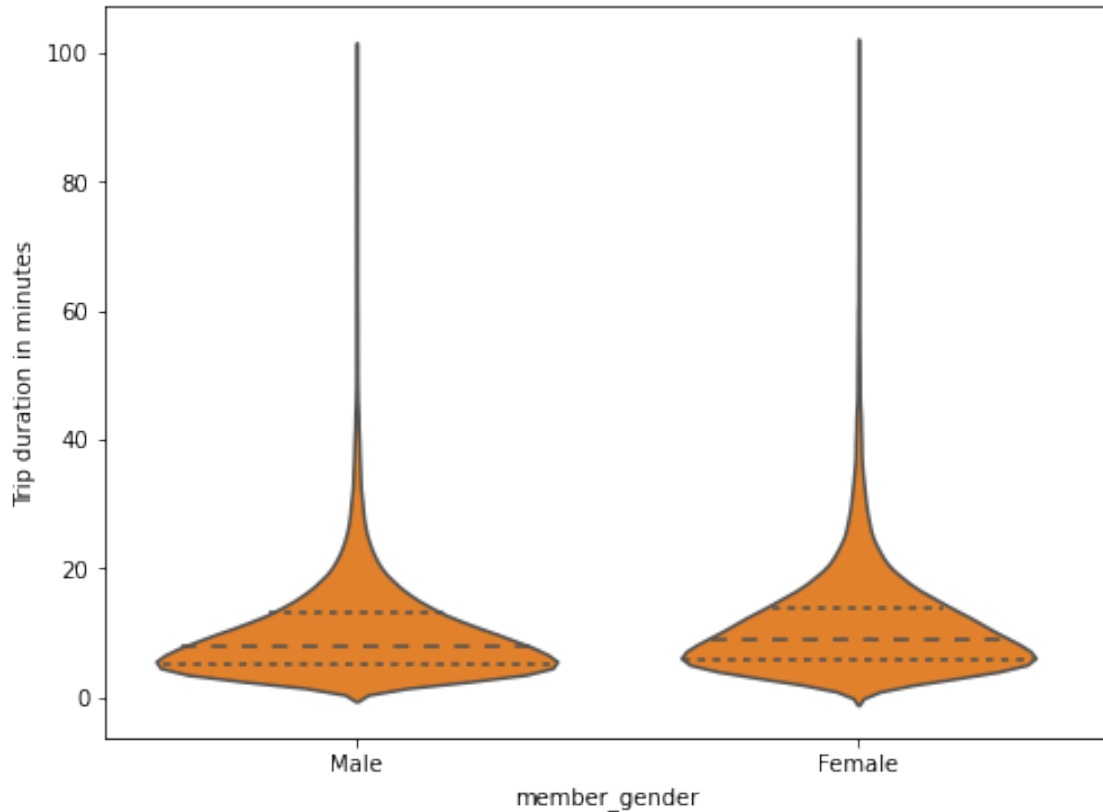
From the graph below we can see that customers overall ride for a longer time - we will take a closer look later at this relationship.

```
[194]: plt.figure(figsize = [8, 6])
sb.boxplot(data=df_clean, x='user_type', y='duration_minutes', color=sb.
↪color_palette()[0]);
plt.xlabel('user_type');
plt.ylabel('Trip duration in minutes');
```




We can also see that Female users are going for longer rides on average

```
[195]: plt.figure(figsize = [8, 6])
sb.violinplot(data=df_clean, x='member_gender', y='duration_minutes', color=sb.
    ↳color_palette()[1], inner='quartile');
plt.xlabel('member_gender');
plt.ylabel('Trip duration in minutes');
```




The hour at which the bikes are rented is also interesting. We can see that people mostly use the bikes around 8 AM and 5 PM. Which indicates that they are used for commuting.

```
[196]: fig = plt.figure(figsize = [16,4])
sb.countplot(data = df_clean, x = 'hour', hue = 'member_gender');
```




```
[197]: fig = plt.figure(figsize = [16,4])
sb.countplot(data = df_clean, x = 'hour', hue = 'user_type');
```




As we can see in the graph below bike renting patterns differ a lot based on whether it's the weekend or a weekday. On weekdays the bikes are used mostly during rush hours. On the weekend, on the other hand, usage is more spread out with a peak during the midday.

```
[198]: fig = plt.figure(figsize = [16,4])
sb.countplot(data = df_clean.query('weekday == ["Monday", "Tuesday",
↪ "Wednesday", "Thursday", "Friday"]'), x = 'hour', hue = 'user_type');
```




```
[199]: fig = plt.figure(figsize = [16,4])
sb.countplot(data = df_clean.query('weekday == ["Saturday", "Sunday"]'), x =
↪ 'hour', hue = 'user_type');
```



```
[200]: days = [ 'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday' ]
        days_dtype = pd.api.types.CategoricalDtype(categories=days, ordered=True)
        df_clean['weekday'] = df_clean['weekday'].astype(days_dtype)
```

Below I wanted to look at the relationship between user type and the day on which the bike is rented. Unfortunately due to the much higher number of subscribers vs customers, it's hard to make any conclusions. Let's try to transform the data.

```
[201]: plt.figure(figsize = [8, 6])
        sb.countplot(data=df_clean, x='weekday', hue='user_type');
        plt.xlabel('weekday');
        plt.ylabel('count');
```



```
[202]: df_clean.user_type.value_counts()
```

```
[202]: Subscriber    153739
        Customer      15774
        Name: user_type, dtype: int64
```

I will create a new data frame where I will calculate a ratio for subscribers and customer's which will help us analyze the data better.


```
[203]: df_ratio = df_clean.groupby(['weekday', 'user_type'])['duration_sec'].count()
df_ratio = df_ratio.reset_index()
```

```
[204]: df_ratio.rename(columns={'duration_sec': 'quantity'}, inplace=True)
```

```
[205]: df_ratio
```

```
[205]:
```

	weekday	user_type	quantity
0	Monday	Customer	2216
1	Monday	Subscriber	22619
2	Tuesday	Customer	2197
3	Tuesday	Subscriber	27563
4	Wednesday	Customer	1997
5	Wednesday	Subscriber	25590
6	Thursday	Customer	2729
7	Thursday	Subscriber	29934
8	Friday	Customer	2418
9	Friday	Subscriber	24374
10	Saturday	Customer	2067
11	Saturday	Subscriber	11864
12	Sunday	Customer	2150
13	Sunday	Subscriber	11795

```
[206]: df_ratio_cust = df_ratio.query('user_type == "Customer"')
df_ratio_cust
```

```
[206]:
```

	weekday	user_type	quantity
0	Monday	Customer	2216
2	Tuesday	Customer	2197
4	Wednesday	Customer	1997
6	Thursday	Customer	2729
8	Friday	Customer	2418
10	Saturday	Customer	2067
12	Sunday	Customer	2150

```
[207]: df_ratio_sub = df_ratio.query('user_type == "Subscriber"')
df_ratio_sub
```

```
[207]:
```

	weekday	user_type	quantity
1	Monday	Subscriber	22619
3	Tuesday	Subscriber	27563
5	Wednesday	Subscriber	25590
7	Thursday	Subscriber	29934
9	Friday	Subscriber	24374
11	Saturday	Subscriber	11864
13	Sunday	Subscriber	11795

```
[208]: df_ratio_cust = df_ratio_cust.assign(ratio=df_ratio_cust.quantity.  
      ↪transform(lambda x: x / x.sum()))
```

```
[209]: df_ratio_sub = df_ratio_sub.assign(ratio=df_ratio_sub.quantity.transform(lambda x:  
      ↪x: x / x.sum()))
```

```
[210]: df_ratio_sub
```

```
[210]:
```

	weekday	user_type	quantity	ratio
1	Monday	Subscriber	22619	0.147126
3	Tuesday	Subscriber	27563	0.179284
5	Wednesday	Subscriber	25590	0.166451
7	Thursday	Subscriber	29934	0.194707
9	Friday	Subscriber	24374	0.158541
11	Saturday	Subscriber	11864	0.077170
13	Sunday	Subscriber	11795	0.076721

```
[211]: ratios = df_ratio_cust.append(df_ratio_sub)
```

```
[212]: ratios
```

```
[212]:
```


	weekday	user_type	quantity	ratio
0	Monday	Customer	2216	0.140484
2	Tuesday	Customer	2197	0.139280
4	Wednesday	Customer	1997	0.126601
6	Thursday	Customer	2729	0.173006
8	Friday	Customer	2418	0.153290
10	Saturday	Customer	2067	0.131038
12	Sunday	Customer	2150	0.136300
1	Monday	Subscriber	22619	0.147126
3	Tuesday	Subscriber	27563	0.179284
5	Wednesday	Subscriber	25590	0.166451
7	Thursday	Subscriber	29934	0.194707
9	Friday	Subscriber	24374	0.158541
11	Saturday	Subscriber	11864	0.077170
13	Sunday	Subscriber	11795	0.076721

```
[213]: ratios.to_csv(r'ratios.csv', index = False, header=True)
```

After transforming the data we can see an interesting observation. During the working week proportionally subscribers used the service more, however, on the weekend regular customer used it proportionally more often. This means that on the weekend many casual bike riders use the service and people that use it during the week for commuting don't use the service as often.

```
[214]: plt.figure(figsize=(16,6))  
sb.barplot(x='weekday', y='ratio', hue='user_type', data=ratios);  
plt.xlabel('Weekday');
```

```
plt.ylabel('Ratio');
plt.xticks(rotation = 45);
```



1.4.1 Talk about some of the relationships you observed in this part of the investigation. How did the feature(s) of interest vary with other features in the dataset?

From my visual analysis I can conclude that:

- With age people take shorter rides
- Customer rent the bikes for a longer time vs subscribers
- People use mostly bikes for commuting (around 8 AM and 5 PM)

1.4.2 Did you observe any interesting relationships between the other features (not the main feature(s) of interest)?

A very interesting observation is that on weekdays on average subscribers use the bikes more - probably for commuting. However, on the weekend random customers on average use the service more - probably leisurely bike riders :)

1.5 Multivariate Exploration

Create plots of three or more variables to investigate your data even further. Make sure that your investigations are justified, and follow from your work in the previous sections.

```
[215]: df_clean.head(1)
```

```
[215]: duration_sec      start_time      end_time  bike_id \
4      1585 2019-02-28 23:54:18.549 2019-03-01 00:20:44.074    4898

      user_type  member_birth_year  member_gender  bike_share_for_all_trip \
```

4	Subscriber	1974.0	Male	Yes
---	------------	--------	------	-----

	weekday	hour	duration_minutes	age
4	Thursday	23	26	45

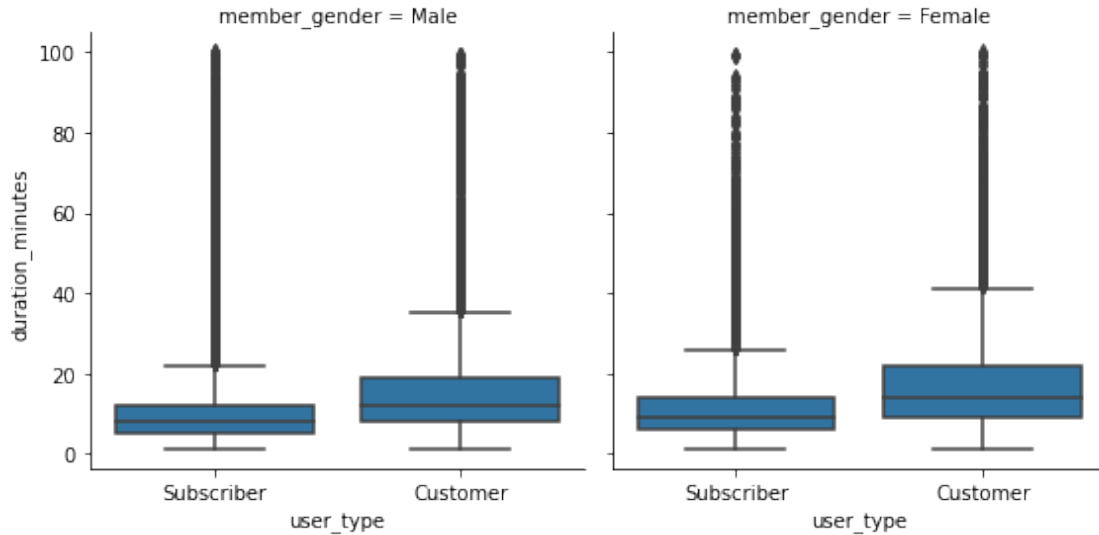
```
[216]: df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 169513 entries, 4 to 183411
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   duration_sec                        169513 non-null  int64
1   start_time                         169513 non-null  datetime64[ns]
2   end_time                           169513 non-null  datetime64[ns]
3   bike_id                            169513 non-null  int64
4   user_type                          169513 non-null  object
5   member_birth_year                  169513 non-null  float64
6   member_gender                      169513 non-null  object
7   bike_share_for_all_trip            169513 non-null  object
8   weekday                            169513 non-null  category
9   hour                               169513 non-null  int64
10  duration_minutes                    169513 non-null  int32
11  age                                 169513 non-null  int32
dtypes: category(1), datetime64[ns](2), float64(1), int32(2), int64(3),
object(3)
memory usage: 18.4+ MB
```

When analyzing the graph below we can see no clear differences in bike usage depending on user type and gender. The relationship is similar - customers be it male or female on average rent the bike for longer periods


```
[217]: g = sb.FacetGrid(data = df_clean, col = 'member_gender', size = 4)
g.map(sb.boxplot, 'user_type', 'duration_minutes')
```

```
[217]: <seaborn.axisgrid.FacetGrid at 0x249689ab160>
```




Below we can see an interesting observation. We can verify that female bike riders tend to rent the bike for longer. Also, we can verify that the youngest people rent the bike for the longer periods of time. However, one interesting thing to note is that people that are older than 50 years tend to rent the bike for longer compared to 30 - 50 year olds. We did not see this when analyzing the graph comparing age and rental duration.

```
[218]: fig = plt.figure(figsize = [16,6])
sb.pointplot(data=df_clean, x='age', y='duration_minutes', hue='member_gender',
             ↪dodge=0.3, linestyle="");
plt.xlabel('age');
plt.ylabel('duration_minute');
```



The graph below also confirms our previous findings. Customers on average rent the bike for longer periods. The youngest people rent also for the longest duration. However here we can also see the confirmation that bike riders over 50 years old start to ride for longer periods.

```
[219]: fig = plt.figure(figsize = [16,6])
sb.pointplot(data=df_clean, x='age', y='duration_minutes', hue='user_type',
             ↪dodge=0.3, linestyle="");
plt.xlabel('age');
plt.ylabel('duration_minute');
```



1.5.1 Talk about some of the relationships you observed in this part of the investigation. Were there features that strengthened each other in terms of looking at your feature(s) of interest?

Most of our previous findings were strengthened in the Multivariate exploration.

- Female riders rent bikes for longer periods
- The youngest people rent a bike for the longest
- Customer rent bikes for a longer duration vs subscribers

1.5.2 Were there any interesting or surprising interactions between features?

The most surprising feature was the age analysis. As previously stated I confirmed that 'The youngest people rent a bike for the longest'. However, what we did not see on our scatter plot is that after the age of 50 people start going for longer rides which is very interesting. It may be connected with a fact that the bike stops being a tool for commuting and starts being a tool for exercising. Mostly we can see this trend when looking at customers. Which as we previously assumed use bikes mostly on weekends for leisure.

All in all this was a very interesting analysis

At the end of your report, make sure that you export the notebook as an html file from the File > Download as... > HTML menu. Make sure you keep track of where the

exported file goes, so you can put it in the same folder as this notebook for project submission. Also, make sure you remove all of the quote-formatted guide notes like this one before you finish your report!

```
[220]: df_clean.to_csv(r'df_clean.csv', index = False, header=True)
```